


Files



- ..
- sample_data
 - Mall_Customers (1).csv
 - Mall_Customers_Clustered.csv
 - heart_disease_uci.csv
 - submission.csv
 - test.csv
 - test1.csv
 - train.csv
 - train1.csv

[6]
✓ 0s

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# 1. Load Kaggle dataset
data = pd.read_csv('train.csv')

# 2. Drop rows where our target or feature might be missing
data = data.dropna(subset=['GrLivArea', 'SalePrice'])

# 3. Select feature and target
X = data[['GrLivArea']]
y = data['SalePrice']

# 4. Split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# 5. Train model
model = LinearRegression()
model.fit(X_train, y_train)

# 6. Predictions
y_pred = model.predict(X_test)

# 7. Evaluation
print(f"R2 Score: {r2_score(y_test, y_pred):.4f}")
print(f"MSE: {mean_squared_error(y_test, y_pred):.2f}")

print("-" * 30)

# 8. Predict a new house (Passing as DataFrame to prevent warnings)
new_house = pd.DataFrame([[2000]], columns=['GrLivArea'])
predicted_price = model.predict(new_house)
print(f"Predicted Price for 2,000 sq ft: ${predicted_price[0]:.2f}")
```

R2 Score: 0.5543

Disk 65.52 GB available

Variables Terminal

gh UV
ow



Search



DELL